

CSC 460
Project 2 Report
March 9th, 2015

Paul Hunter
Justin Guze

Table Of Contents

<u>1.0 Overview</u>
<u>2.0 RTOS Environment</u>
<u>2.1 C Run Time</u>
<u>2.2 RTOS Data Structures</u>
<u>2.2.1 Task Descriptors and Periodic Meta Data</u>
<u>2.2.2 Task Queues</u>
<u>2.2.3 Context Switching</u>
<u>3.0 OS API</u>
<u>4.0 Task Scheduling and OS Requests</u>
<u>4.1 Task Life Cycle</u>
<u>4.2 Kernel Request Handling and Dispatch</u>
<u>4.3 Scheduling</u>
<u>5.0 Services</u>
<u>5.1 Initialization</u>
<u>5.2 Subscribing</u>
<u>5.3 Publishing</u>
<u>6.0 Error Handling</u>
<u>7.0 Testing</u>
<u>7.1 Test Cases</u>
<u>8.0 Profiling</u>
<u>9.0 Conclusion</u>

1.0 Overview

For this project we designed and implemented a simple real-time operating system to fulfill the specifications provided to us. This report will describe the design, implementation, testing, and profiling of this active kernel style system, as well as any insights or recommendations we have. The target hardware platform for this project was again the ATmega2560.

Within this simple operating system, there exists three types of tasks, System, Periodic, and Round Robin. Systems tasks are of the highest priority, execute on a first come first serve basis, and always run to completion before yielding. Periodic tasks, which are the second highest priority, are time critical tasks which are executed based on a information provided by the engineer, defined by a start tick, period, and worst case execution time. Finally, Round Robin tasks are the lowest priority, are first come first serve, but only execute for a single tick before yielding the processor to the next Round Robin task.

One of the key requirements of the system was to ensure preemption of tasks by any tasks with a higher priority. For example, if a Round Robin task were to create a System level task, as will be discussed later in detail, the creation of the task would result in the Kernel switching to execute this new System task, before returning to execute the Round Robin task.

In addition to prioritized and periodic tasks, the OS also supports a simple synchronization mechanism known as a Service. Services allow a non-periodic task to subscribe to the publishing of a value by another task, the task which subscribes will wait until another task publishes a value to that Service. This mechanism will be discussed in detail below in Section 5.

In the sections which follow we will discuss the Runtime environment of the OS and the data structures it uses, the life cycle of tasks and how they are scheduled, our testing strategies and test cases, and finally, a profile of the execution times and overhead of the OS. Lastly, we will present any notes or recommendations we have having completed this project.

2.0 RTOS Environment

The sections which follow provide detailed descriptions of the runtime environment and memory space of the operation system. We designed the system to contain only statically allocated data within the kernel, although this requires definition of the limits of stack and number of tasks at compile time, it ensures that memory usage by the kernel is consistent from run to run, regardless of when during execution tasks are created and destroyed.

2.1 C Run Time

Although we were provided with examples of c runtime instructions to initialize the system, we chose to use the included crt.s that comes bundled with AVR. The decision was made to keep the implementation simple and more portable as the initialization is written in C; execution starts with the OS' main method, which then calls kernel_init to initialize the needed data structures.

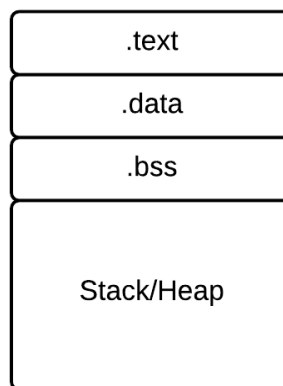


Figure 1 - High Level Memory allocation after C Runtime Initialization

The .text portion of memory contains all program code, including the Kernel and User Application. The .data portion of memory contains all variables declared and initialized, either by the user application or the kernel. For the kernel, task contexts including stacks, periodic task metadata, current task descriptors are stored here, frankly, most of the data related to the kernel can be found in this section. The .bss section contains all variables declared but not initialized, again either by the user application or our kernel, as a result, it contains the remaining data of the kernel, such as the current task pointer, and kernel stack pointer value.

The remaining free space in memory is then used for the Heap and Stack. Heap is used for dynamic allocation of resources, the kernel, as it is strictly statically defined, will never make use of this space, however an ill-informed user may write an application using this space, often we try to steer clear of dynamic memory in embedded systems however. The Stack space is then normally used as the stack of the process, however, applications written using our OS, will not make use of this space, as its stack space for each task is defined in static memory space in the .data section. If a developer

were to find there was excessive free space in the stack, one may choose to increase the workspace of the tasks.

2.2 RTOS Data Structures

The subsections below discuss in detail the data structures and allocation schemes of contexts and task descriptors.

2.2.1 Task Descriptors and Periodic Meta Data

Tasks within the system are defined using the `task_descriptor_t` type, shown below, which contains all components to a tasks execution but the code itself, which as mentioned previously is located in `.text`.

```
typedef struct td_struct
{
    uint8_t stack[WORKSPACE];
    uint8_t* volatile sp;
    task_priority_t priority;
    periodic_task_metadata_t* periodic_desc;
    task_state_t state;
    int arg;
    int16_t * data;
    struct td_struct* next;
} task_descriptor_t;
```

As one would expect the we find the task's stack, stack pointer, priority, and state (more on States in section 2.1.2). In addition to these intuitive elements, task descriptors also contain a pointer to a `periodic_task_metadata_t` structure, a pointer to an integer which is used to publish values to the task if it has subscribed, and finally, a pointer to another `task_descriptor_t`, details of these three additional fields will be discussed below.

Unlike System and Round Robin tasks, Periodic tasks have additional information that defines them; this information is stored in a `periodic_task_metadata_t`, shown below, which contains the period, worst case execution time, and next, a value which indicates the next tick the task is to be executed. In addition, the structure also contains a pointer to a task descriptor, as well as another periodic metadata structure. The use of these fields will be discussed below.

```

struct ptd_metadata_struct
{
    struct td_struct* task;
    uint16_t period;
    uint16_t wcet;
    uint16_t next;
    ptd_metadata_struct* nextT;
} ;
typedef ptd_metadata_struct periodic_task_metadata_t;

```

2.2.2 Task Queues

To store and queue tasks for execution, the kernel makes use of 5 queues. To avoid the use of dynamic memory these queues and the maximum number of tasks and periodic tasks are defined at compile time and allocated statically. Previously we mentioned the that periodic metadata and task descriptors each contain pointers to themselves and the other. This allows us to create queues in place, without the need to allocate nodes as the descriptors themselves are the nodes in the linked list.

At startup, the kernel initializes all task descriptors and periodic metadata in place, and creates two 'dead' queues which contain all the unused descriptors for tasks and periodic metadata structures. As tasks are created, takes are moved to an appropriate queue based on their priority. As tasks are killed, their task descriptors are cleared for reuse, and the descriptor is added back to the dead pool.

Below in Figure 2, we have defined the system to have a maximum of 4 tasks, up to 2 of which can be periodic. As a result, we statically allocate an array of 4 task_descriptor_t structures, and another array of 2 periodic_metadata_t. After initialization, the structures are all linked in order, and the queue for the dead tasks points to the first and last entry.

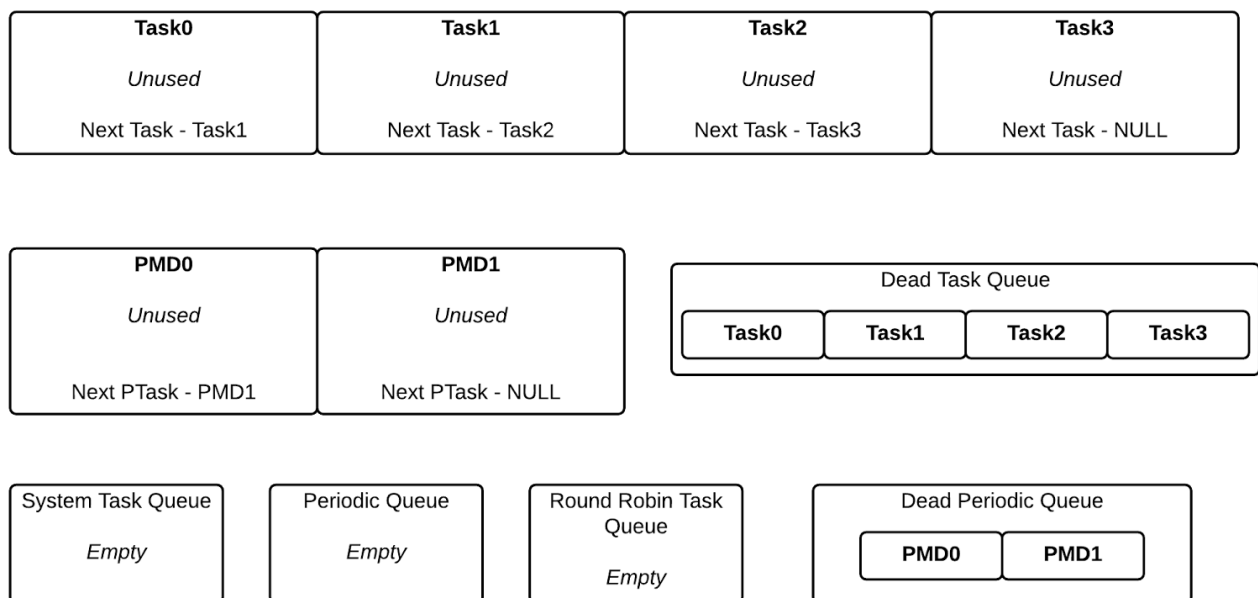


Figure 2 - Task Descriptors and Periodic Meta Data Arrays and Queues after Initialization

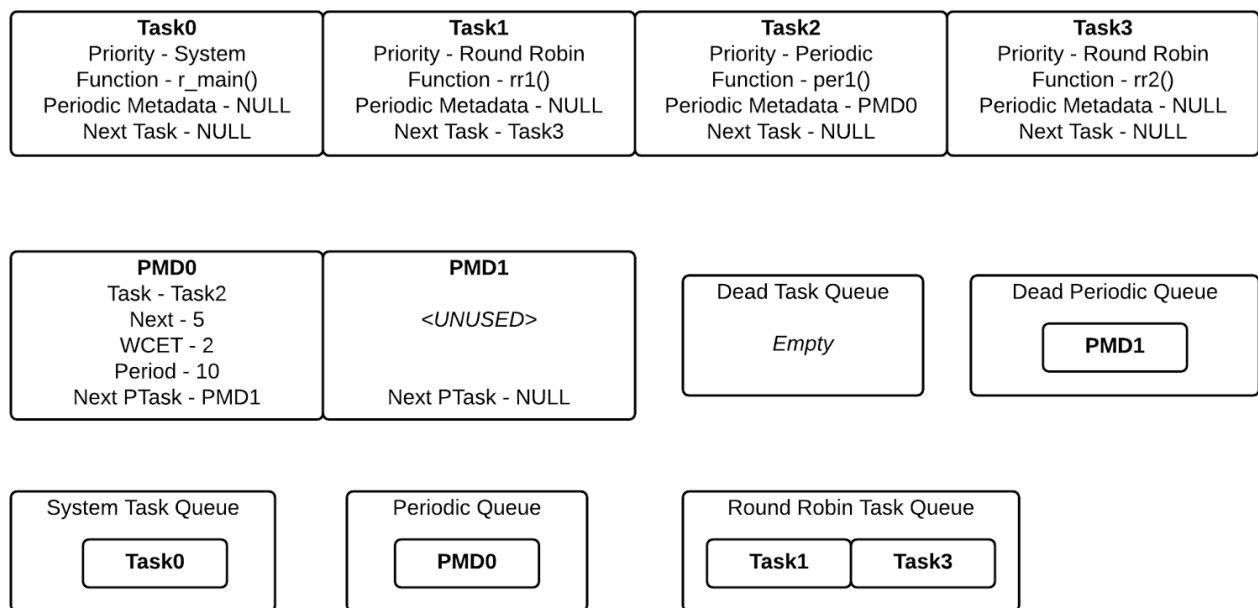


Figure 3 - Task Descriptors and Periodic Metadata Arrays and Queues after task creation

In Figure 3 above, the system has created a system task, two round robin tasks, and a single periodic task. As we can see, the system task and round robin tasks have been moved from the dead queue and into their respective schedule queue. The single periodic task however, is not placed in a queue using its task descriptor, instead it is placed in a queue based on its metadata, which contains a reference to the task descriptor for use by the scheduler. Figure 3 below shows the state of the queues after the creation of these tasks.

Section 4.0 discusses task scheduling in detail and provides additional details on how queues are utilized for scheduling and task management.

2.2.3 Context Switching

As we decided to implement the OS using an active kernel, API calls require a switch between User space and Kernel space, and this requires a change of execution context. In addition, this context switch is needed to change back to the executing task once the kernel has completed its work. The context switch operation requires that you replace the current execution context of the running task, with that of the kernel or a user task.

The execution context of a task includes its stack pointer, all 32 general purpose registers, the SREG register and the EIND value, which stores an additional bit used in function addresses in combination with the two ZH and ZL registers. To accomplish the swap of these values with that of the target task, we made use of the provided RESTORE_CTX() and SAVE_CTX() macros, with a modification to include the missing EIND value, required. This additional bit also changes the size of the stack pointer and return address, to compensate, we had to add two additional bytes to the initial stack of the task when constructing its first execution context.

Switching occurs by first pushing the executing context of the current task onto the task's stack, the execution of the target task, or kernel, is then restored by popping it off its stack, where it would have been stored when changing tasks. This approach requires that we must manually build the first execution context by hand when creating a task.

3.0 OS API

Detailed documentation for the OS interface can be found in os.h. Below in Figure 4 we show the available functions of the kernel. Many of the calls require a context change to the kernel, but calls to Task_GetArg() and Now() do not, instead they can fetch directly into memory for the values they are looking for.

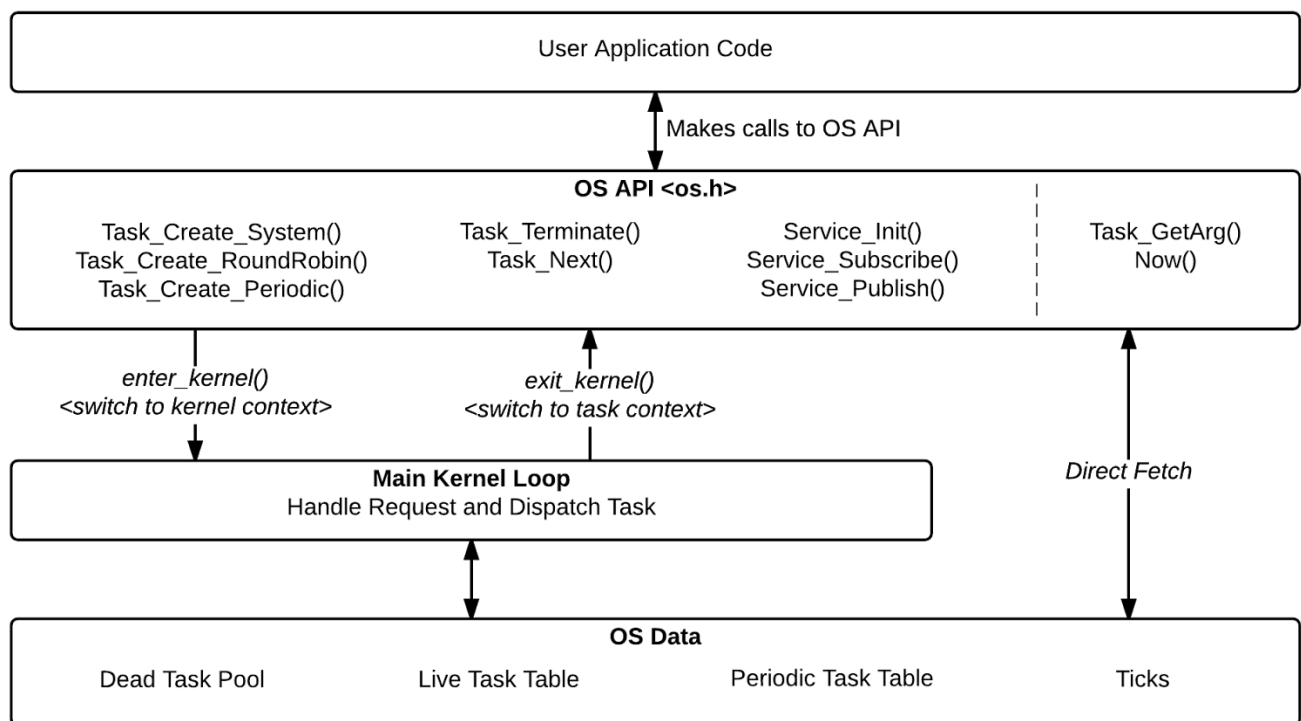


Figure 4 - OS API and Context Boundaries

4.0 Task Scheduling and OS Requests

After the system is initialized, a single System level task is created before the OS starts, that task points to the user space `r_main`, defined by the user application. This `r_main` should then create tasks as needed and otherwise configure the application on the user side. We chose to run this `r_main` as a System task to ensure that it is not preempted by a task created during user space initialization. This section will cover in detail the life cycle of a task as well as scheduling and preemption schemes implemented in our RTOS.

4.1 Task Life Cycle

Tasks have a finite number of potential states which they can occupy. Then not in user, we refer to the task as being DEAD. After initial creation, and when the thread is not waiting, it resides in the READY state, indicating it can be run at any time. The current running task is the sole task with the state RUNNING, while any task waiting on a service will reside in the WAITING state.

Below in Figure 5, the life cycle of a task is given as a finite state machine. Each of the operations used as transition triggers is described briefly below, more detailed information on services can be found below in Section 5.

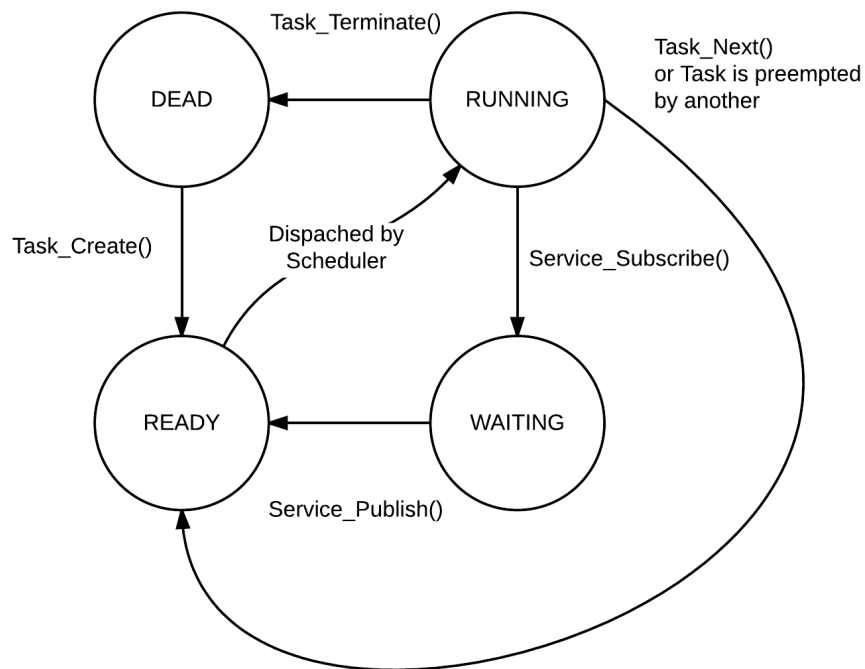


Figure 5 - Task Life cycle

- *Task_Next* - Task voluntarily relinquishes its remaining time.
- *Task_Create* - Task creates a new task of any priority. It is possible that a call to this OS method, for example by a Round Robin task to create a System task will result in immediate preemption, placing the current task back into a ready state, and in this case, placing the Round Robin task at the front of the Round Robin Queue.
- *Task_Terminate* - This is called automatically by tasks which simply return from their functions as a result of the way we constructed their initial execution context, but can also be called explicitly to terminate the task, clearing its descriptor and periodic metadata, it is then added to the dead queues.
- *Service_Subscribe* - The task will wait on a service to publish before returning to a READY state. Note, periodic tasks may not call this method as they may not be bound to synchronization mechanisms, as per the RTOS specification.
- *Service_Publish* - Task publishes a value to the service, waking any tasks currently waiting on that service. This call can result in preemption if a task of high priority is awoken.

4.2 Kernel Request Handling and Dispatch

With each OS API call, interrupts are disabled, the current stack register is saved, and we enter the kernel via `enter_kernel`, which switches context to that of kernel space. By always coming back to Kernel space on every request which could potentially trigger a preemption, the kernel is able to seamlessly select the highest priority task and run it, having already effectively suspended the last task during the context switch it can place it back in its appropriate queue.

4.3 Scheduling

Although the Periodic Process Plan (PPP) is an effective way to plan period tasks, we did not find it all that intuitive. As a result, we replaced it with a dynamically scheduled system, unfortunately this means the system is not able to detect time conflicts until they occur. The schedule is created by sorting the `periodic_metadata_t` by their `now` field, indicating the next tick for execution.

As mentioned previously, each call to the OS causes a context change, this allows us to reschedule in the moment after the request has been handled. Each request handler is responsible for changing the state of the current task to ready if it is to be preempted. If the handler does change the state, for example of a Round Robin task created a new System task, `kernel_dispatch` will select a new task to run based on its priority.

If in the event that no tasks are ready to run, either because there are none or they are all waiting for services to publish, or are not yet scheduled to run, an Idle Task is used. This idle task is created as an additional task with its own context, and resides in the last index of the `task_descriptor_t` array. This task, which is essentially an endless loop, allows us to keep the processor busy waiting for interrupts which may cause another task to become ready, and is never placed inside a queue.

5.0 Services

As specified in `os.h`, services are used to transfer data safely between tasks, as well as allow tasks to wait for certain data to be published. This has many uses, such as waiting for a section of memory to become available, or waiting for radio communication to finish with a device. As per the documentation, services were split into three distinct OS API calls: `Service_Init`, `Service_Subscribe`, and `Service_Publish`.

The service data structure is built up of only a task queue, the same sort of queues used for system tasks and round robin tasks. This queue simply holds all tasks that subscribe to a given service.

```
struct service
{
    task_queue_t task_queue;
};
```

Each OS API function ends up making a call to the kernel. As seen in Figure 3, all Service functions create a new kernel request to process, enter the kernel, and allow the kernel to handle all appropriate service functionality. We decided to put all service functionality at kernel level instead of user level because all task manipulation happens at kernel level, regardless of the overhead added when switching contexts. It also had the benefit of easily being able to preempt tasks that awoke when publishing data to a service. The following sections will go into more detail describing each call.

5.1 Initialization

As described in the given documentation, `Service_Init` will return a pointer to a new service type. First, the user makes a call to the OS API `Service_Init`, which will disable interrupts and set the kernel request to 'SERVICE_INIT'. It then switches context and enters kernel mode, where the kernel will proceed to create a new service descriptor using an unused one, found in the services array in the OS data. It will then exit the kernel, and return the service descriptor to the user. There is a maximum amount of services that can be created in the RTOS, which in this case is 10. If a user creates more than the maximum amount, the RTOS will abort using error code `ERR_RUN_8_SERVICE_CAPACITY_REACHED`.

5.2 Subscribing

A task may provide a service descriptor and the address to where the task expects the publisher to publish the data. Once the task subscribes, it will enter a waiting state until a task publishes to the service. A task will call the OS API `Service_Subscribe`, which will proceed to create a kernel request called `SERVICE_SUB`. Then, the task will switch contexts to the kernel, which will subscribe the specified task to the service descriptor. The task is added to the service's task queue, and its state is

changed to WAITING. Then, the kernel will attempt to dispatch a new task, since the current task is set to WAITING.

The Subscription queues use the same mechanism the task queues for scheduling do, utilizing the pointer field within the task descriptor. As a result, keeping this queue also requires no additional memory after compile time.

There is a restriction on the type of tasks that can subscribe to a service. Only system and round robin tasks may subscribe. If a periodic task attempts to subscribe to a service, the RTOS will abort and send the `ERR_RUN_10_PERIODIC_SUBSCRIBE` to the user. If the user provides an invalid service descriptor, the RTOS will abort and send the appropriate `ERR_RUN_9_INVALID_SERVICE`.

5.3 Publishing

As described in the documentation, `Service_Publish` will take a service descriptor and an `int16_t` to publish to every task found in the services task queue. All tasks found in the service change their state to READY and moved from the task queue of the service to its appropriate scheduling queue; the data is written to the expected address given when the tasks subscribed. More specifically, a call to the OS API `Service_Publish` will create a kernel request called `SERVICE_PUB`. It will then switch contexts to the kernel, and publish the data to the given service descriptor. Then, for each task in the service's task queue, we will set their state to READY, and add them back into the appropriate task queue depending on their priority (SYSTEM or ROUND_ROBIN). It will then check and see if there are any tasks that preempt the currently running task, and if so, will change the current task accordingly.

If an invalid service descriptor is provided to the `Service_Publish` function, the RTOS will abort and report the `ERR_RUN_9_INVALID_SERVICE` error message to the user using LEDs.

There were various problems attempting to implement the publish routine. Initially, every publish call would go to completion, even if there was a system task that was awoken when publishing to the service. This was quickly fixed by adding the tasks back into their appropriate queues. Another problem we found was handling interrupts that published to services. The desired functionality was to prevent preemption on an interrupt, and let them run to completion.

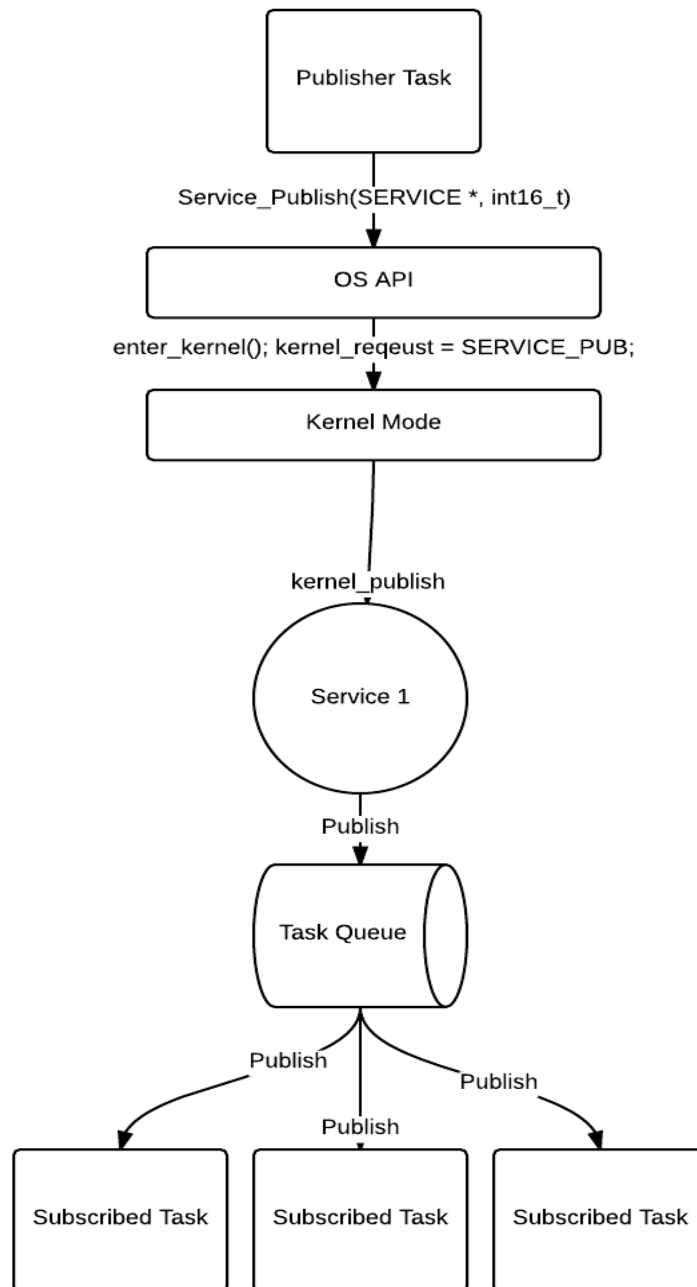


Figure 6 - Publisher Subscriber Call Diagram

6.0 Error Handling

The sample RTOS code came bundled with a number of runtime and initializations errors that could be encountered by the system. We were fans of the 'pulse x times to indicate error code x' and extended it into our solution. When an error occurs, we set the error_msg variable to the appropriate error code, and call OS_Abort(), disabling interrupts and entering a loop, stopping the execution of kernel and all other tasks. We removed and added various error codes, which are listed below as well as in error_code.h.

ERR_1_INIT_FAILURE

An initialization error happened during the function kernel_init when first starting the RTOS.

ERR_RUN_0_USER_CALLED_OS_ABORT

A runtime error meant to catch any instance of the user calling OS_Abort without actually setting the error message. In a production environment, this error should never happen.

ERR_RUN_1_TOO_MANY_TASKS

A runtime error stating the user is attempting to create too many tasks. The maximum allowed tasks is currently set in the MAXPROCESS constant.

ERR_RUN_2_TOO_MANY_PERIODIC_TASKS

A runtime error that catches if the user is creating specifically too many periodic tasks. A second constant defines the number of periodic tasks, MAXPROCESSPER.

ERR_RUN_3_PERIODIC_INVALID_CONFIGURATION

The user specified a configuration that cannot work for periodic tasks. Specifically, this error will be thrown when the user specifies a period of 0 or they set the worst case execution time to be longer than the period of the task.

ERR_RUN_4_PERIODIC_TOOK_TOO_LONG

A runtime error found when a periodic task runs longer than its expected worst case execution time.

ERR_RUN_5_PERIODIC_TASKS_SCHEDULED_AST

A runtime error that occurs when the user schedule results in two periodic tasks being scheduled to execute at the same time. It is a requirement that the engineer ensures this never happens, as the RTOS system does not provide failsafes.

ERR_RUN_6_ILLEGAL_ISR_KERNEL_REQUEST

A runtime error when an interrupt service routine attempts to make a kernel request that is not valid, for example, Terminate_Task() or Service_Subscribe()

ERR_RUN_7_RTOS_INTERNAL_ERROR

An error which is not well defined occurred during runtime. Further investigation would be needed.

ERR_RUN_8_SERVICE_CAPACITY_REACHED

A runtime error found when a user attempts to create more services than the max amount of services specified under the constant MAXSERVICES.

ERR_RUN_9_INVALID_SERVICE

A runtime error found when the user gives the functions Service_Subscribe or Service_Publish an invalid service descriptor.

ERR_RUN_10_PERIODIC_SUBSCRIBE

A runtime error thrown when a periodic task attempts to subscribe to a service. There is a strict requirement that subscribing tasks must not be periodic.

ERR_RUN_11_PERIODIC_FOUND_WHEN_PUBLISHING

A runtime error thrown when a periodic task is awoken when a service is published to. Since services don't accept periodic subscribers, there is no way this task should happen unless the user misuses the service descriptor without running the OS API functions.

ERR_RUN_12_TASK_WITHOUT_PRIORITY

A runtime error found when a task that is awoken from a published service has no priority specified. This should never happen unless the user misuses the service descriptor.

7.0 Testing

Rigorous testing was applied to the RTOS to ensure we caught every possible task state and combination, as well as testing all our possible error states. The testing framework used is taken mostly courtesy of the test framework used by Justin Tanner and Scott Craig during their initial creation of the RTOS.

The test framework uses `trace.h` and `uart.h` to print messages to the trace buffer, and output it to the serial monitor. To allow us to view the trace messages, we ended up opening the Arduino IDE and checking that the messages found in the serial monitor were identical to the expected messages found in each test file.

Every test is a specific file found in the tests folder of the project. Each test has a description in its filename to specify the type of test that is being run, and there is a comment at the top of each test displaying the expected string to be found in the serial monitor. Since every test should run with a newly started RTOS, it was necessary for each test to define its own `r_main`. Then, a constant was used to specify which test's `r_main` was compiled. This meant that each tested needed us to recompile the entire RTOS, and flash it onto the boards memory every time.

Each test was built around the same concept: Write a test to produce a specific functionality, and use trace to ensure that everything ran in the appropriate order. Generally we ended up writing the numbers 1 - 10 to trace and ensuring they appeared in chronological order. For tests that had multiple tasks and services, it was beneficial to write chronological trace statements to prove the order of initialization and completion.

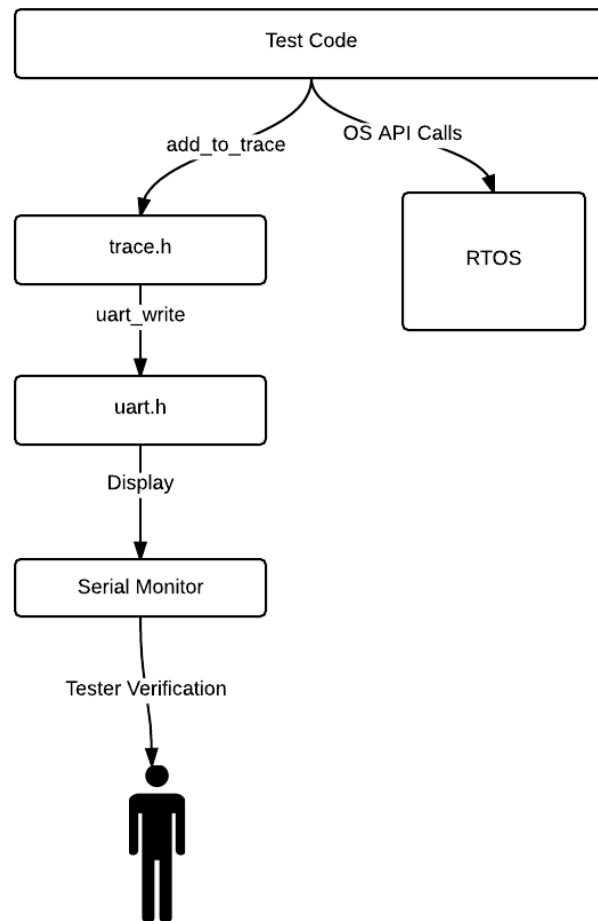


Figure 8 - Test Framework Call Diagram

7.1 Test Cases

The following is a list of all test cases run on the RTOS, and a short description of their use case.

test000_sanity.cpp

The basic sanity check test. It ensures we can create a system task that will write a message to the trace buffer. If we can't see any output from this, then either our test framework or our RTOS is broken and all other tests will provide useless information.

test001_NowCorrectTime.cpp

Ensures that the Now function returns the correct time in ms. This test will create a loop that writes the result from Now to the trace buffer, and the waits 100ms. Visually, we inspect and ensure the values being written are in increments of 100.

test002_create_system_task.cpp

Creates a system task, and ensures that it is scheduled and runs to completion. This is the basic system task test. If this fails, then all subsequent tests that use system tasks should fail as well.

test003_multiple_system_task_order.cpp

Creates multiple systems tasks to be scheduled. Ensures that system tasks run in the order they are created.

test004_system_task_preemption.cpp

Creates two system tasks, and ensures that one system task doesn't preempt the other. System tasks should run in the order they are scheduled. Very similar to test003.

test005_system_preempt_nonsystem_task.cpp

Creates a round robin task that then creates a system task. The system task should preempt the round robin task as it has more priority over it. This is verified using the order of the trace messages that is written by each task. Ensures our basic preemption code is working correctly.

test006_create_periodic_task.cpp

Creates a periodic task that should get scheduled and will then print a message to the trace buffer. This is the basic periodic task test. If this test fails, then all subsequent periodic tests will also fail.

test007_periodic_worst_case_time_larger_period.cpp

This test will create a periodic task that has a worst case execution time larger than its period. We expect this test to throw the error ERR_RUN_3_INVALID_CONFIGURATION and flash the red LED on the board 3 times.

test008_periodic_zero_period.cpp

This test will create a periodic task where the period is specified to be 0. We expect this test to throw the error ERR_RUN_3_INVALID_CONFIGURATION and flash the red LED on the board 3 times.

test009_multiple_periodic_tests.cpp

Creates multiple periodic tasks that do not collide, and ensures they are both running at the correct times. This is verified by printing both the TICK they are running, as well as a trace message corresponding to that task. This test ensures we are scheduling periodic tasks correctly, and that we aren't forgetting to reschedule periodic tasks when they call Task_Next.

test010_periodic_tasks_collide.cpp

This test will create two periodic tasks that start at the exact same time. We expect this test to throw the error ERR_RUN_5_PERIODIC_TASKS_SCHEDULED_AST and flash the red LED on the board 5 times.

test011_periodic_task_not_completed_on_time.cpp

This test will create a periodic task with a period of 5 ms. The task itself is specified to delay for 50 ms. We expect this test to throw the error ERR_RUN_4_PERIODIC_TOOK_TOO_LONG and flash the red LED 4 times. This test ensures our logic that checks the amount of time a periodic task has run for is working as expected.

test012_periodic_system_preemption.cpp

This test will create a periodic task that will then create a system task. The periodic task should then be preempted by the system task. This is verified by the order in which they write messages to the trace buffer. This test ensures that our logic for system preemption works also for periodic tasks.

test013_create_round_robin.cpp

This test creates a round robin test, and ensures it gets scheduled and runs to completion. This is the basic round robin test case, and if it fails then every other test that uses round robin tasks will also fail.

test014_multiple_round_robin.cpp

This test will create multiple round robin tasks, and ensure they are scheduled and run in the correct order by viewing the order in which they wrote to the trace buffer.

test015_create_service.cpp

This test will call the Service_Init function, and ensure we return a service descriptor. This test is to make sure our service initialization calls aren't throwing errors for any reason.

test016_create_max_services.cpp

This test ensures we can create up to MAXSERVICES. It will loop, and call Service_Init MAXSERVICES times.

test017_create_over_max_services.cpp

This test will attempt to create more than MAXSERVICES services. We expect this test to throw the error ERR_RUN_8_SERVICE_CAPACITY_REACHED and blink the red LED on the board 8 times.

test018_system_subscribe_service.cpp

This test ensures that a system task can subscribe to a service. A system task is created, that will then subscribe to a service. We ensure that the task is waiting by checking the trace messages and ensuring the system task trace message was not written to the buffer.

test019_periodic_subscribe_service.cpp

This test will attempt to create a periodic task that will then subscribe to a service. As periodic tasks should not be able to subscribe, we expect this test to throw the error ERR_RUN_10_PERIODIC_SUBSCRIBE and blink the red LED 10 times.

test020_round_robin_subscribe_service.cpp

This test will ensure a round robin task can subscribe to a service, and will then wait. It will create a round robin task that will immediately subscribe to a service. We ensure it is waiting by verifying that it did not write to the trace buffer.

test021_publish_round_robin.cpp

This test verifies that round robin tasks can publish to a service and wake up other round robin tasks. We also ensure the order they are scheduled in is correct.

test022_publish_preempt_system.cpp

This test verifies that system tasks awoken from a service preempt the running round robin task. A round robin task will publish to a service that contains a system task. This system task should then

preempt the round robin task that published, and run to completion before running the round robin task.

test023_subscribe_invalid_service.cpp

This test will attempt to subscribe to a service using an invalid service descriptor. We expect it to throw the error `ERR_RUN_9_INVALID_SERVICE` and blink the red LED on the board 9 times.

test024_service_interrupted.cpp

This test ensures that an interrupt service routine does not get preempted when publishing to a service. If an ISR publishes to a service that contains a system task, we expect the ISR to complete its task before switching contexts to the system task.

test025_publish_periodic.cpp

This test ensures that a periodic task can publish to a service. It will also ensure that tasks will be rescheduled in the appropriate order.

test026_rr_task_next.cpp

This test ensures that a round robin task can call `Task_Next` and switch contexts properly.

8.0 Profiling

Profiling was run on the RTOS to measure the overhead added by the new OS API functions. As shown in Figure 3, every new function added except the Now function switches contexts into the kernel to complete the request. It then becomes necessary for us to measure the timing issues this could cause.

Most of the profiling code was taken courtesy of Mike Krazanowski. The original source code can be found [here](#) and the report detailing his profile findings can be found [here](#).

To measure the amount of time each new function took, we included various macros to send a HIGH or LOW signal to specific pins. To test any function, right before calling it we would send a HIGH signal to a pin. Then, once the function completed, we would send the LOW signal to the pin. By using a logic analyzer, we were able to measure how long a certain pin produced a HIGH signal.

Function Call	Time Elapsed (microseconds)
OS_Init	47.1975
Task_Create_System	43.9375
Task_Create_Periodic	52.6875
Task_Create_RoundRobin	43.8125
Task_Next	33.8125
Service_Init	26.8750
Service_Subscribe	34.9375
Service_Publish	44.8750

Table 1 - Profile data for the new function calls

OS_Init

The system initialization took 47.1975 microseconds to complete. This is mainly the busywork needed to initialize the data in each task queue, ensuring linked lists are empty, setting up times, and creating the initial r_main system task.

Task_Create_System

This function took 43.9375 microseconds to complete. This function requires a change in context from user mode to kernel mode. From there, the kernel creates a new system task and adds it to the back of the system queue. It will then dispatch the next highest priority task, and exit the kernel, switching contexts again from kernel mode to user mode. Overall, having two context switches for task creation does not produce a significant amount of overhead.

Task_Create_RoundRobin

This function took 43.8125 microseconds to complete. It runs in a nearly identical fashion to the Task_Create_System, which is why their elapsed time is also nearly identical. It similarly has two context switches in and out of kernel mode.

Task_Create_Periodic

This function took 52.6875 microseconds to complete. Although very similar to Task_Create_RoundRobin and Task_Create_System, the extra time on this function is most likely due to the priority queue used for periodic tasks, as well as the overhead necessary for computing the next time they should be run.

Task_Next

This function took 33.8125 microseconds to complete. This task sets a kernel request, and switches contexts from user mode to kernel mode. It will then relinquish control of the current task, and find the next highest priority task to complete. The main time constraint of this function is the context switch into kernel mode and back.

Service_Init

This function took 26.8750 microseconds to complete. This function is one of the simplest functions in the system. It switches contexts from kernel mode to user mode, and pulls an unused service descriptor from the services list. The overhead of this task is mainly the context switching that occurs, and could actually all be done in user mode instead.

Service_Subscribe

This function took 34.9375 microseconds to complete. This function makes a context switch into kernel mode to add tasks to the service. It was a little difficult to test specifically a task subscribe function, since the subscribing task would get set to WAITING after subscribing. To alleviate this, we queued up another low priority task that would send a LOW signal to the pin. Overall, the overhead of this task is mainly the context switch into kernel mode, along with the time it takes to queue up the task into the service.

Service_Publish

This function took 44.8750 microseconds to complete. This task is dependant on the number of tasks already in the service queue, as it takes time to dequeue a task and publish the data to the expected address. In this case, we had one task in the service queue when publishing. The extra time taken to complete this task is due to the extra scheduling logic added to this task. We must both enqueue the tasks from the service back into their respective ready queues, as well as ensure our current task did not get preempted by any task that awoke.

Overall, the overhead added by context switching into kernel mode does not currently seem to play a critical role in timing. It seems to take a significant time dequeuing tasks from our queues and service queues, which may pose a problem in the future when we attempt to scale up the amount of tasks in the system.

9.0 Conclusion

Overall, we managed to complete the RTOS specification according to the given documentation, as well as fully test and profile our code. We created a fully functioning real time operating system that can create 3 different kinds of tasks (System, Periodic, Round Robin) as well as initialize services, subscribe to services, and publish to services. We managed to work around task preemption that occurs in various edge cases, and we proved during our testing that the scheduling works at the appropriate ticks and in the appropriate order.

We did run into a few issues attempting to implement this RTOS. We initially had a test case that always failed when running multiple periodic tests, where the second task created never ran to completion. Upon a code inspection, it was discovered that the linked list that handled periodic tasks was implemented incorrectly. Once fixed, the test worked as expected.

Another problem we ran into was initially getting the RTOS code onto the ATmega2560 board. It took a surprising amount of time to install the correct tools to easily push working code onto the board.

We had a lot of trouble handling the specific case where interrupts would attempt to publish to a service. Since we allow preemption of tasks to happen during publishing, we would always allow a system task to preempt the publishing task. Unfortunately, we did not desire preemption to happen when an interrupt was publishing, because we always wanted the interrupt to run to completion first.

In the end, the RTOS works as specified, and will work as long as the user creates tasks and services according to the specification. We look forward to modifying the RTOS code to work efficiently for Project 3.